

# LAB 3.2

## Introduction To Pandas Module

Here, we are learning data analysis. Pandas is a software library written for the Python programming language that helps with data manipulation and analysis. It offers various data structures which help in the manipulation of data. One of the data structures used is DataFrame. DataFrame is used as a way to store data in a rectangular grid. Similar to how data is stored in Microsoft Excel.

## Installation

We can easily install this library using the terminal. Follow the steps given below to install Pandas: Step 1: Go to the Windows PowerShell. Step 2: We can install pandas with the help of pip easily. Head over to the directory where you want to install the pandas and fire the below command: **pip install pandas** This is very simple to install.

```
# !pip install pandas
```

## Other way to install

If the above steps do not work for you, here is another way to install: Step 1: Go to your browser. Search for **python unofficial binaries**, or click <https://www.lfd.uci.edu/~gohlke/pythonlibs/> and you will be redirected to the binary page. Step 2: To download the library, search for pandas – Ctrl + F will help you. Download pandas library based on your python version. Step 3: After the installation, head over to the directory where you've downloaded the pandas. Step 4: Right-click to open Windows PowerShell. Pip can help us with the installation. Just go to that location in your terminal and run the following command: **pip install pandas** Complete this command by the TAB button so that the location of your folder is mentioned.

## import pandas

Once Pandas is installed, import it in your applications by adding the import keyword: **import pandas as pd** Now Pandas is imported and ready to use.

```
import pandas

dict1 = {
    'Name': ["Umar", "Ali", "Usman", "Samra", "behzad"],
    'Marks': [80, 72, 77, 65, 15],
    'Class': ["1st Year", "2nd Year", "1st Year", "3rd Year", "2nd Year"]
}

df = pandas.DataFrame(dict1)
df = pandas.DataFrame(dict1)
```

```
df
```

	Name	Marks	Class
0	Umar	80	1st Year
1	Ali	72	2nd Year
2	Usman	77	1st Year
3	Samra	65	3rd Year
4	behzad	15	2nd Year

## Pandas as pd

Pandas is usually imported under the pd alias.

alias:

In Python alias are an alternate name for referring to the same thing. Create an alias with the as keyword while importing: **import pandas as pd** Now the Pandas package can be referred to as pd instead of pandas.

```
import pandas as pd
```

```
df = pd.DataFrame(dict1)
```

```
df
```

	Name	Marks	Class
0	Umar	80	1st Year
1	Ali	72	2nd Year
2	Usman	77	1st Year
3	Samra	65	3rd Year
4	behzad	15	2nd Year

## Checking Pandas Version

The version string is stored under **version** attribute.

```
pd.__version__
```

```
'2.2.3'
```

## Import to CSV

```
df.to_csv('pythonMonday.csv')
```

## Remove Index

```
df.to_csv('pythonMondayRemoveIndex.csv', index=False)
```

## Show some rows from the Start

```
df.head()
```

	Name	Marks	Class
0	Umar	80	1st Year
1	Ali	72	2nd Year
2	Usman	77	1st Year
3	Samra	65	3rd Year
4	behzad	15	2nd Year

```
df.head()
```

	Name	Marks	Class
0	Umar	80	1st Year
1	Ali	72	2nd Year
2	Usman	77	1st Year
3	Samra	65	3rd Year
4	behzad	15	2nd Year

```
df.tail(3)
```

	Name	Marks	Class
2	Usman	77	1st Year
3	Samra	65	3rd Year
4	behzad	15	2nd Year

```
df.tail()
```

	Name	Marks	Class
0	Umar	80	1st Year
1	Ali	72	2nd Year
2	Usman	77	1st Year
3	Samra	65	3rd Year
4	behzad	15	2nd Year

```
df.describe()
```

	Marks
count	5.000000
mean	61.800000
std	26.771253
min	15.000000
25%	65.000000
50%	72.000000
75%	77.000000
max	80.000000

```
classData= pd.read_csv("dataset.csv")
```

```
classData
```

	Unnamed: 0	name	Marks	Class	gpa	class
0	0	Umar	80	1st Year	4.0	1
1	1	Behzad	72	2nd Year	3.0	2
2	2	Ali	77	1st Year	3.5	4
3	3	Samra	65	3rd Year	2.4	6
4	4	Usman	15	2nd Year	2.0	5

```
classData.drop(0, axis=0)
```

	Unnamed: 0	name	Marks	Class	gpa	class
1	1	Behzad	72	2nd Year	3.0	2
2	2	Ali	77	1st Year	3.5	4
3	3	Samra	65	3rd Year	2.4	6
4	4	Usman	15	2nd Year	2.0	5

```
classData
```

	Unnamed: 0	name	Marks	Class	gpa	class
0	0	Umar	80	1st Year	4.0	1
1	1	Behzad	72	2nd Year	3.0	2
2	2	Ali	77	1st Year	3.5	4
3	3	Samra	65	3rd Year	2.4	6
4	4	Usman	15	2nd Year	2.0	5

```
classData['gpa']
```

0	4.0
1	3.0
2	3.5
3	2.4
4	2.0

```
Name: gpa, dtype: float64
```

```
classData['Class']
```

0	1st Year
1	2nd Year
2	1st Year
3	3rd Year
4	2nd Year

```
Name: Class, dtype: object
```

```
classData['class']
```

0	1
1	2
2	4
3	6
4	5

```
Name: class, dtype: int64
```

```
classData['name']
```

```
0    Umar
1    Behzad
2    Ali
3    Samra
4    Usman
Name: name, dtype: object
```

```
type(classData['name'])
```

```
pandas.core.series.Series
```

```
classData['gpa']
```

```
0    4.0
1    3.0
2    3.5
3    2.4
4    2.0
Name: gpa, dtype: float64
```

```
classData['gpa'][0]
```

```
np.float64(4.0)
```

```
classData['gpa'][0]=90
```

```
C:\Users\PMLS\AppData\Local\Temp\ipykernel_100812\3783213339.py:1:
```

```
SettingWithCopyWarning:
```

```
A value is trying to be set on a copy of a slice from a DataFrame
```

```
See the caveats in the documentation:
```

```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#  
returning-a-view-versus-a-copy
```

```
classData['gpa'][0]=90
```

```
classData
```

	Unnamed: 0	name	Marks	Class	gpa	class
0	0	Umar	80	1st Year	90.0	1
1	1	Behzad	72	2nd Year	3.0	2
2	2	Ali	77	1st Year	3.5	4
3	3	Samra	65	3rd Year	2.4	6
4	4	Usman	15	2nd Year	2.0	5

```
classData['gpa'][4]
```

```
np.float64(2.0)
```

```
classData['gpa'][4]=5
```

```
C:\Users\PMLS\AppData\Local\Temp\ipykernel_100812\3957951123.py:1:
```

```
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
classData['gpa'][4]=5
```

```
classData["Class"][4]="Second Year"
```

C:\Users\PMLS\AppData\Local\Temp\ipykernel\_100812\3056981296.py:1:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
classData["Class"][4]="Second Year"
```

```
classData
```

	Unnamed: 0	name	Marks	Class	gpa	class
0	0	Umar	80	1st Year	90.0	1
1	1	Behzad	72	2nd Year	3.0	2
2	2	Ali	77	1st Year	3.5	4
3	3	Samra	65	3rd Year	2.4	6
4	4	Usman	15	Second Year	5.0	5

```
classData["name"][4]
```

```
'Usman'
```

```
classData
```

	Unnamed: 0	name	Marks	Class	gpa	class
0	0	Umar	80	1st Year	90.0	1
1	1	Behzad	72	2nd Year	3.0	2
2	2	Ali	77	1st Year	3.5	4
3	3	Samra	65	3rd Year	2.4	6
4	4	Usman	15	Second Year	5.0	5

```
classData.to_csv('pythonMondayNew.csv')
```

```
classData.to_csv('pythonMondayNewRemoveIndex.csv', index=False)
```

## Create Labels

With the index argument, you can name your own labels.

```
classData.index = ["First", "Second", "Third", "Fourth", "Fifth"]
```

```
classData
```

	Unnamed: 0	name	Marks	Class	gpa	class
First	0	Umar	80	1st Year	90.0	1
Second	1	Behzad	72	2nd Year	3.0	2
Third	2	Ali	77	1st Year	3.5	4
Fourth	3	Samra	65	3rd Year	2.4	6
Fifth	4	Usman	15	Second Year	5.0	5

```
classData["name"]["Fifth"]
```

```
'Usman'
```

## Pandas

Pandas is an open source data analysis library written in Python. It uses the power and speed of numpy to make data analysis and preprocessing easy for data scientists. **Pandas** is used to analyze data. Python data analysis using pandas is a great start to Exploratory Data Analysis (EDA) for a data scientist.

## Pandas Data Structure

Panda has two types of Data Structure: a) Series b) Dataframe

### Series

It is a one dimensional array with indexes. It stores a single column or row of data in a Dataframe. It is a one dimensional array capable of holding any type of data.

### Dataframe

It is a tabular spreadsheet like structure representing rows each of which contains one or more columns. It is a multi dimensional array capable of holding any type of data. **Series** is like a column, a **DataFrame** is the whole table.

```
classData['gpa']
```

```
First    90.0
```

```
Second   3.0
```

```
Third    3.5
```

```
Fourth   2.4
```

```
Fifth    5.0
```

```
Name: gpa, dtype: float64
```

```
type(classData['gpa'])
```

```
pandas.core.series.Series
```

```
classData
```

	Unnamed: 0	name	Marks	Class	gpa	class
First	0	Umar	80	1st Year	90.0	1

Second	1	Behzad	72	2nd Year	3.0	2
Third	2	Ali	77	1st Year	3.5	4
Fourth	3	Samra	65	3rd Year	2.4	6
Fifth	4	Usman	15	Second Year	5.0	5

```
type(classData)
```

```
pandas.core.frame.DataFrame
```

```
import numpy as np
```

```
series1= pd.Series(np.random.rand())
```

```
series1
```

```
0    0.863799
dtype: float64
```

```
series1= pd.Series(np.random.rand(30))
```

```
series1
```

```
0    0.490505
1    0.642388
2    0.701803
3    0.427315
4    0.162072
5    0.823410
6    0.007146
7    0.124420
8    0.964552
9    0.691485
10   0.833831
11   0.320095
12   0.537346
13   0.236504
14   0.427526
15   0.726029
16   0.171906
17   0.505061
18   0.558176
19   0.255227
20   0.468784
21   0.236258
22   0.372229
23   0.948222
24   0.285313
25   0.570974
26   0.910419
27   0.877554
28   0.276910
```



```
29      0.403935
dtype: float64
```

```
type(series1)
```

```
pandas.core.series.Series
```

```
newDf = pd.DataFrame(np.random.rand(330,5))
# newDf = pd.DataFrame(np.random.rand(330,5), index=np.arange(330))
```

```
type(newDf)
```

```
pandas.core.frame.DataFrame
```

```
newDf
```

	0	1	2	3	4
0	0.428644	0.845682	0.785568	0.887098	0.655208
1	0.771389	0.351744	0.090406	0.778027	0.017863
2	0.563619	0.783972	0.415338	0.971640	0.643814
3	0.975824	0.777449	0.590080	0.534420	0.627735
4	0.525620	0.740227	0.516741	0.533727	0.041541
...	...	...	...	...	...
325	0.439123	0.909931	0.741812	0.729477	0.929148
326	0.156356	0.658465	0.396061	0.002193	0.819785
327	0.163087	0.522659	0.529201	0.883081	0.600300
328	0.776183	0.177136	0.876736	0.473567	0.575241
329	0.478343	0.311120	0.342621	0.845378	0.528896

```
[330 rows x 5 columns]
```

```
newDf.head(2)
```

	0	1	2	3	4
0	0.428644	0.845682	0.785568	0.887098	0.655208
1	0.771389	0.351744	0.090406	0.778027	0.017863

```
newDf.dtypes
```

```
0      float64
1      float64
2      float64
3      float64
4      float64
dtype: object
```

```
newDf.tail(2)
```

	0	1	2	3	4
328	0.776183	0.177136	0.876736	0.473567	0.575241
329	0.478343	0.311120	0.342621	0.845378	0.528896

```
newDf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 330 entries, 0 to 329
```

```
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
0	0	330 non-null	float64
1	1	330 non-null	float64
2	2	330 non-null	float64
3	3	330 non-null	float64
4	4	330 non-null	float64

```
dtypes: float64(5)
```

```
memory usage: 13.0 KB
```

```
newDf.describe()
```

	0	1	2	3	4
count	330.000000	330.000000	330.000000	330.000000	330.000000
mean	0.508833	0.508141	0.499055	0.496970	0.486661
std	0.292255	0.272950	0.280110	0.289980	0.275773
min	0.001118	0.000351	0.009838	0.000110	0.006137
25%	0.254701	0.289593	0.251430	0.268755	0.259932
50%	0.518867	0.529780	0.497992	0.482537	0.491128
75%	0.762419	0.734844	0.733904	0.771237	0.700446
max	0.999580	0.990092	0.993643	0.999924	0.999810

```
newDf.dtypes
```

```
0    float64
```

```
1    float64
```

```
2    float64
```

```
3    float64
```

```
4    float64
```

```
dtype: object
```

```
newDf.iloc[0, 0] = "Umar"
```

```
C:\Users\PMLS\AppData\Local\Temp\ipykernel_100812\392238210.py:1:
```

```
FutureWarning: Setting an item of incompatible dtype is deprecated and  
will raise an error in a future version of pandas. Value 'Umar' has  
dtype incompatible with float64, please explicitly cast to a  
compatible dtype first.
```

```
newDf.iloc[0, 0] = "Umar"
```

```
newDf.dtypes
```

```
0    object
```

```
1    float64
```

```
2    float64
```

```
3    float64
```

```
4      float64
dtype: object
```

```
newDf
```

	0	1	2	3	4
0	Umar	0.845682	0.785568	0.887098	0.655208
1	0.771389	0.351744	0.090406	0.778027	0.017863
2	0.563619	0.783972	0.415338	0.971640	0.643814
3	0.975824	0.777449	0.590080	0.534420	0.627735
4	0.52562	0.740227	0.516741	0.533727	0.041541
...	...	...	...	...	...
325	0.439123	0.909931	0.741812	0.729477	0.929148
326	0.156356	0.658465	0.396061	0.002193	0.819785
327	0.163087	0.522659	0.529201	0.883081	0.600300
328	0.776183	0.177136	0.876736	0.473567	0.575241
329	0.478343	0.311120	0.342621	0.845378	0.528896

```
[330 rows x 5 columns]
```

```
newDf.index
```

```
RangeIndex(start=0, stop=330, step=1)
```

```
newDf.columns
```

```
RangeIndex(start=0, stop=5, step=1)
```

```
newDf.to_numpy()
```

```
array([[ 'Umar', 0.8456821947437894, 0.7855679557200089,
         0.8870980510759304, 0.6552082421872513],
       [0.7713890033519027, 0.3517441707532226, 0.0904061315213649,
         0.7780271654897173, 0.01786332949894054],
       [0.5636187515945378, 0.7839715463074424, 0.41533796014304836,
         0.9716402421640051, 0.6438144530947115],
       ...,
       [0.16308672038126282, 0.5226594156849904, 0.5292009521323162,
         0.8830814913507018, 0.6002999227158293],
       [0.776183422835709, 0.177136421257725, 0.8767359753400374,
         0.4735673253418551, 0.5752408473426548],
       [0.47834250340156226, 0.31112029439200506, 0.34262138437345857,
         0.8453784051793725, 0.5288955586446525]], dtype=object)
```

```
df.loc[0, 'Name'] = 'Umar'
```

```
newDf.head()
```

	0	1	2	3	4
0	Umar	0.845682	0.785568	0.887098	0.655208
1	0.771389	0.351744	0.090406	0.778027	0.017863
2	0.563619	0.783972	0.415338	0.971640	0.643814

```
3 0.975824 0.777449 0.590080 0.534420 0.627735
4 0.52562 0.740227 0.516741 0.533727 0.041541
```

```
newDf.T
```

```
      0      1      2      3      4      5      6
\
0      Umar 0.771389 0.563619 0.975824 0.52562 0.93062
0.008113
1 0.845682 0.351744 0.783972 0.777449 0.740227 0.135287
0.715199
2 0.785568 0.090406 0.415338 0.59008 0.516741 0.478412
0.219151
3 0.887098 0.778027 0.97164 0.53442 0.533727 0.656286
0.452222
4 0.655208 0.017863 0.643814 0.627735 0.041541 0.255964
0.229701

      7      8      9      ...      320      321      322
323 \
0 0.145869 0.351346 0.570656 ... 0.165955 0.621539 0.650114
0.886066
1 0.667175 0.295342 0.738414 ... 0.246964 0.307061 0.523363
0.735252
2 0.616049 0.115213 0.361927 ... 0.575118 0.660315 0.356277
0.193276
3 0.507092 0.313166 0.977326 ... 0.903633 0.546056 0.048854
0.212792
4 0.695681 0.692118 0.203833 ... 0.992035 0.860566 0.097351
0.600351

      324      325      326      327      328      329
0 0.327249 0.439123 0.156356 0.163087 0.776183 0.478343
1 0.8256 0.909931 0.658465 0.522659 0.177136 0.31112
2 0.173012 0.741812 0.396061 0.529201 0.876736 0.342621
3 0.285449 0.729477 0.002193 0.883081 0.473567 0.845378
4 0.456484 0.929148 0.819785 0.6003 0.575241 0.528896
```

```
[5 rows x 330 columns]
```

```
newDf.sort_index(axis=0)
```

```
      0      1      2      3      4
0      Umar 0.845682 0.785568 0.887098 0.655208
1 0.771389 0.351744 0.090406 0.778027 0.017863
2 0.563619 0.783972 0.415338 0.971640 0.643814
3 0.975824 0.777449 0.590080 0.534420 0.627735
4 0.52562 0.740227 0.516741 0.533727 0.041541
...
325 0.439123 0.909931 0.741812 0.729477 0.929148
```

326	0.156356	0.658465	0.396061	0.002193	0.819785
327	0.163087	0.522659	0.529201	0.883081	0.600300
328	0.776183	0.177136	0.876736	0.473567	0.575241
329	0.478343	0.311120	0.342621	0.845378	0.528896

[330 rows x 5 columns]

newDf.sort\_index(axis=0, ascending=False)

	0	1	2	3	4
329	0.478343	0.311120	0.342621	0.845378	0.528896
328	0.776183	0.177136	0.876736	0.473567	0.575241
327	0.163087	0.522659	0.529201	0.883081	0.600300
326	0.156356	0.658465	0.396061	0.002193	0.819785
325	0.439123	0.909931	0.741812	0.729477	0.929148
...	...	...	...	...	...
4	0.52562	0.740227	0.516741	0.533727	0.041541
3	0.975824	0.777449	0.590080	0.534420	0.627735
2	0.563619	0.783972	0.415338	0.971640	0.643814
1	0.771389	0.351744	0.090406	0.778027	0.017863
0	Umar	0.845682	0.785568	0.887098	0.655208

[330 rows x 5 columns]

newDf.sort\_index(axis=1)

	0	1	2	3	4
0	Umar	0.845682	0.785568	0.887098	0.655208
1	0.771389	0.351744	0.090406	0.778027	0.017863
2	0.563619	0.783972	0.415338	0.971640	0.643814
3	0.975824	0.777449	0.590080	0.534420	0.627735
4	0.52562	0.740227	0.516741	0.533727	0.041541
...	...	...	...	...	...
325	0.439123	0.909931	0.741812	0.729477	0.929148
326	0.156356	0.658465	0.396061	0.002193	0.819785
327	0.163087	0.522659	0.529201	0.883081	0.600300
328	0.776183	0.177136	0.876736	0.473567	0.575241
329	0.478343	0.311120	0.342621	0.845378	0.528896

[330 rows x 5 columns]

newDf.sort\_index(axis=1, ascending=False)

	4	3	2	1	0
0	0.655208	0.887098	0.785568	0.845682	Umar
1	0.017863	0.778027	0.090406	0.351744	0.771389
2	0.643814	0.971640	0.415338	0.783972	0.563619
3	0.627735	0.534420	0.590080	0.777449	0.975824
4	0.041541	0.533727	0.516741	0.740227	0.52562
...	...	...	...	...	...
325	0.929148	0.729477	0.741812	0.909931	0.439123

```

326  0.819785  0.002193  0.396061  0.658465  0.156356
327  0.600300  0.883081  0.529201  0.522659  0.163087
328  0.575241  0.473567  0.876736  0.177136  0.776183
329  0.528896  0.845378  0.342621  0.311120  0.478343

```

```
[330 rows x 5 columns]
```

```
newDf.head()
```

```

      0      1      2      3      4
0    Umar  0.845682  0.785568  0.887098  0.655208
1  0.771389  0.351744  0.090406  0.778027  0.017863
2  0.563619  0.783972  0.415338  0.971640  0.643814
3  0.975824  0.777449  0.590080  0.534420  0.627735
4  0.52562  0.740227  0.516741  0.533727  0.041541

```

```
newDf.iloc[1, 0]
```

```
0.7713890033519027
```

```
type(newDf.iloc[:, 0])
```

```
pandas.core.series.Series
```

```
newDf
```

```

      0      1      2      3      4
0    Umar  0.845682  0.785568  0.887098  0.655208
1  0.771389  0.351744  0.090406  0.778027  0.017863
2  0.563619  0.783972  0.415338  0.971640  0.643814
3  0.975824  0.777449  0.590080  0.534420  0.627735
4  0.52562  0.740227  0.516741  0.533727  0.041541
..    ..    ..    ..    ..    ..
325  0.439123  0.909931  0.741812  0.729477  0.929148
326  0.156356  0.658465  0.396061  0.002193  0.819785
327  0.163087  0.522659  0.529201  0.883081  0.600300
328  0.776183  0.177136  0.876736  0.473567  0.575241
329  0.478343  0.311120  0.342621  0.845378  0.528896

```

```
[330 rows x 5 columns]
```

## view

```
newDfView=newDf
```

```
newDfView
```

```

      0      1      2      3      4
0    Umar  0.845682  0.785568  0.887098  0.655208
1  0.771389  0.351744  0.090406  0.778027  0.017863
2  0.563619  0.783972  0.415338  0.971640  0.643814

```

```

3    0.975824  0.777449  0.590080  0.534420  0.627735
4    0.52562  0.740227  0.516741  0.533727  0.041541
...
325  0.439123  0.909931  0.741812  0.729477  0.929148
326  0.156356  0.658465  0.396061  0.002193  0.819785
327  0.163087  0.522659  0.529201  0.883081  0.600300
328  0.776183  0.177136  0.876736  0.473567  0.575241
329  0.478343  0.311120  0.342621  0.845378  0.528896

```

[330 rows x 5 columns]

```
newDfView[0][0]=765
```

C:\Users\PMLS\AppData\Local\Temp\ipykernel\_100812\4063526219.py:1:  
FutureWarning: ChainedAssignmentError: behaviour will change in pandas 3.0!

You are setting values through chained assignment. Currently this works in certain cases, but when using Copy-on-Write (which will become the default behaviour in pandas 3.0) this will never work to update the original DataFrame or Series, because the intermediate object on which we are setting values will behave as a copy. A typical example is when you are setting values in a column of a DataFrame, like:

```
df["col"][row_indexer] = value
```

Use `df.loc[row\_indexer, "col"] = values` instead, to perform the assignment in a single step and ensure this keeps updating the original `df`.

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
newDfView[0][0]=765
```

C:\Users\PMLS\AppData\Local\Temp\ipykernel\_100812\4063526219.py:1:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
newDfView[0][0]=765
```

```
newDf.head()
```

```

      0      1      2      3      4
0    765  0.845682  0.785568  0.887098  0.655208
1  0.771389  0.351744  0.090406  0.778027  0.017863
2  0.563619  0.783972  0.415338  0.971640  0.643814

```

```
3  0.975824  0.777449  0.590080  0.534420  0.627735
4   0.52562  0.740227  0.516741  0.533727  0.041541
```

```
newDfView.head()
```

```
      0      1      2      3      4
0    765  0.845682  0.785568  0.887098  0.655208
1  0.771389  0.351744  0.090406  0.778027  0.017863
2  0.563619  0.783972  0.415338  0.971640  0.643814
3  0.975824  0.777449  0.590080  0.534420  0.627735
4   0.52562  0.740227  0.516741  0.533727  0.041541
```

```
newDfCopy=newDf.copy()
```

```
newDfCopy[0][0]=656
```

```
C:\Users\PMLS\AppData\Local\Temp\ipykernel_100812\2532441962.py:1:
FutureWarning: ChainedAssignmentError: behaviour will change in pandas
3.0!
```

You are setting values through chained assignment. Currently this works in certain cases, but when using Copy-on-Write (which will become the default behaviour in pandas 3.0) this will never work to update the original DataFrame or Series, because the intermediate object on which we are setting values will behave as a copy.

A typical example is when you are setting values in a column of a DataFrame, like:

```
df["col"][row_indexer] = value
```

Use `df.loc[row_indexer, "col"] = values` instead, to perform the assignment in a single step and ensure this keeps updating the original `df`.

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
newDfCopy[0][0]=656
```

```
C:\Users\PMLS\AppData\Local\Temp\ipykernel_100812\2532441962.py:1:
```

```
SettingWithCopyWarning:
```

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation:

[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
newDfCopy[0][0]=656
```

```
newDf.head()
```

```
      0      1      2      3      4
0    765  0.845682  0.785568  0.887098  0.655208
```



1	0.771389	0.351744	0.090406	0.778027	0.017863
2	0.563619	0.783972	0.415338	0.971640	0.643814
3	0.975824	0.777449	0.590080	0.534420	0.627735
4	0.52562	0.740227	0.516741	0.533727	0.041541

```
newDfCopy.head()
```

	0	1	2	3	4
0	656	0.845682	0.785568	0.887098	0.655208
1	0.771389	0.351744	0.090406	0.778027	0.017863
2	0.563619	0.783972	0.415338	0.971640	0.643814
3	0.975824	0.777449	0.590080	0.534420	0.627735
4	0.52562	0.740227	0.516741	0.533727	0.041541

## loc

```
newDf.head(2)
```

	0	1	2	3	4
0	765	0.845682	0.785568	0.887098	0.655208
1	0.771389	0.351744	0.090406	0.778027	0.017863

```
newDf.loc[0,0]
```

```
765
```

```
newDf.loc[0,0]=5678
```

```
newDf.head(2)
```

	0	1	2	3	4
0	5678	0.845682	0.785568	0.887098	0.655208
1	0.771389	0.351744	0.090406	0.778027	0.017863

```
newDf.loc[0,4]
```

```
np.float64(0.6552082421872513)
```

```
newDf.loc[325,1]
```

```
np.float64(0.9099311354265363)
```

```
newDf.loc[0,0]=9875
```

```
newDf.head(2)
```

	0	1	2	3	4
0	9875	0.845682	0.785568	0.887098	0.655208
1	0.771389	0.351744	0.090406	0.778027	0.017863

```
newDf.columns=list("ABCDE")
```

```
newDf.head(3)
```

	A	B	C	D	E
0	9875	0.845682	0.785568	0.887098	0.655208
1	0.771389	0.351744	0.090406	0.778027	0.017863
2	0.563619	0.783972	0.415338	0.971640	0.643814

```
newDf.loc[0, "A"]
```

```
9875
```

```
newDf.loc[0, "A"]=65
```

```
newDf.head(3)
```

	A	B	C	D	E
0	65	0.845682	0.785568	0.887098	0.655208
1	0.771389	0.351744	0.090406	0.778027	0.017863
2	0.563619	0.783972	0.415338	0.971640	0.643814

```
newDf.loc[35, "A"]
```

```
0.45685661234086095
```

```
newDf.drop(0, axis=0).head(3)
```

	A	B	C	D	E
1	0.771389	0.351744	0.090406	0.778027	0.017863
2	0.563619	0.783972	0.415338	0.971640	0.643814
3	0.975824	0.777449	0.590080	0.534420	0.627735

```
newDf.drop("A", axis=1).head(3)
```

	B	C	D	E
0	0.845682	0.785568	0.887098	0.655208
1	0.351744	0.090406	0.778027	0.017863
2	0.783972	0.415338	0.971640	0.643814

```
newDf.head(3)
```

	A	B	C	D	E
0	65	0.845682	0.785568	0.887098	0.655208
1	0.771389	0.351744	0.090406	0.778027	0.017863
2	0.563619	0.783972	0.415338	0.971640	0.643814

```
newDf = newDf.drop("A", axis=1)
```

```
newDf.head(3)
```

	B	C	D	E
0	0.845682	0.785568	0.887098	0.655208
1	0.351744	0.090406	0.778027	0.017863
2	0.783972	0.415338	0.971640	0.643814

```
newDf.loc[[1,2], ["C", "D"]]
```

	C	D
1	0.090406	0.778027
2	0.415338	0.971640

```
newDf.head(3)
```

	B	C	D	E
0	0.845682	0.785568	0.887098	0.655208
1	0.351744	0.090406	0.778027	0.017863
2	0.783972	0.415338	0.971640	0.643814

```
newDf.loc[:,["C","D"]]
```

	C	D
0	0.785568	0.887098
1	0.090406	0.778027
2	0.415338	0.971640
3	0.590080	0.534420
4	0.516741	0.533727
...	...	...
325	0.741812	0.729477
326	0.396061	0.002193
327	0.529201	0.883081
328	0.876736	0.473567
329	0.342621	0.845378

```
[330 rows x 2 columns]
```

```
newDf.loc[[1,2],:]
```

	B	C	D	E
1	0.351744	0.090406	0.778027	0.017863
2	0.783972	0.415338	0.971640	0.643814

```
newDf.loc[:5,["C","D"]]
```

	C	D
0	0.785568	0.887098
1	0.090406	0.778027
2	0.415338	0.971640
3	0.590080	0.534420
4	0.516741	0.533727
5	0.478412	0.656286

```
newDf.loc[[1,2],:"D"]
```

	B	C	D
1	0.351744	0.090406	0.778027
2	0.783972	0.415338	0.971640

```
newDf.loc[(newDf["B"]<0.3)]
```

	B	C	D	E
5	0.135287	0.478412	0.656286	0.255964
8	0.295342	0.115213	0.313166	0.692118
10	0.036757	0.993643	0.087791	0.962828
19	0.227323	0.255985	0.701267	0.156683
21	0.156348	0.636665	0.390927	0.391794
...	...	...	...	...
312	0.183028	0.883414	0.775590	0.560528
313	0.011312	0.603112	0.564635	0.743004
315	0.015563	0.031942	0.601035	0.508256
320	0.246964	0.575118	0.903633	0.992035
328	0.177136	0.876736	0.473567	0.575241

[85 rows x 4 columns]

```
newDf.loc[(newDf["B"]<0.3) & (newDf["D"]>0.2)]
```

	B	C	D	E
5	0.135287	0.478412	0.656286	0.255964
8	0.295342	0.115213	0.313166	0.692118
19	0.227323	0.255985	0.701267	0.156683
21	0.156348	0.636665	0.390927	0.391794
23	0.296591	0.962421	0.584525	0.559112
...	...	...	...	...
312	0.183028	0.883414	0.775590	0.560528
313	0.011312	0.603112	0.564635	0.743004
315	0.015563	0.031942	0.601035	0.508256
320	0.246964	0.575118	0.903633	0.992035
328	0.177136	0.876736	0.473567	0.575241

[69 rows x 4 columns]

```
newDf.head(3)
```

	B	C	D	E
0	0.845682	0.785568	0.887098	0.655208
1	0.351744	0.090406	0.778027	0.017863
2	0.783972	0.415338	0.971640	0.643814

## iloc

```
newDf.iloc[0,0]
```

```
np.float64(0.8456821947437894)
```

```
newDf.iloc[0,0]=555
```

```
newDf.head()
```

	B	C	D	E
0	555.000000	0.785568	0.887098	0.655208

1	0.351744	0.090406	0.778027	0.017863
2	0.783972	0.415338	0.971640	0.643814
3	0.777449	0.590080	0.534420	0.627735
4	0.740227	0.516741	0.533727	0.041541

```
newDf.iloc[[0,4],[1,3]]
```

	C	E
0	0.785568	0.655208
4	0.516741	0.041541

```
newDf.head(4)
```

	B	C	D	E
0	555.000000	0.785568	0.887098	0.655208
1	0.351744	0.090406	0.778027	0.017863
2	0.783972	0.415338	0.971640	0.643814
3	0.777449	0.590080	0.534420	0.627735

```
newDf.drop(0).head()
```

	B	C	D	E
1	0.351744	0.090406	0.778027	0.017863
2	0.783972	0.415338	0.971640	0.643814
3	0.777449	0.590080	0.534420	0.627735
4	0.740227	0.516741	0.533727	0.041541
5	0.135287	0.478412	0.656286	0.255964

```
# newDf.drop(0, axis=1)
```

```
# newDf.drop("B", axis=1).head()
```

```
newDf.drop(["C", "D"], axis=1).head()
```

	B	E
0	555.000000	0.655208
1	0.351744	0.017863
2	0.783972	0.643814
3	0.777449	0.627735
4	0.740227	0.041541

```
newDf.head()
```

	B	C	D	E
0	555.000000	0.785568	0.887098	0.655208
1	0.351744	0.090406	0.778027	0.017863
2	0.783972	0.415338	0.971640	0.643814
3	0.777449	0.590080	0.534420	0.627735
4	0.740227	0.516741	0.533727	0.041541

```
newDf.drop(['B', 'D'], axis=1, inplace=True)
```

```
newDf.head()
```

	C	E
0	0.785568	0.655208
1	0.090406	0.017863
2	0.415338	0.643814
3	0.590080	0.627735
4	0.516741	0.041541

```
newDf.reset_index().head()
```

	index	C	E
0	0	0.785568	0.655208
1	1	0.090406	0.017863
2	2	0.415338	0.643814
3	3	0.590080	0.627735
4	4	0.516741	0.041541

```
newDf.reset_index(drop=True).head()
```

	C	E
0	0.785568	0.655208
1	0.090406	0.017863
2	0.415338	0.643814
3	0.590080	0.627735
4	0.516741	0.041541

```
newDf.head()
```

	C	E
0	0.785568	0.655208
1	0.090406	0.017863
2	0.415338	0.643814
3	0.590080	0.627735
4	0.516741	0.041541

```
newDf.reset_index(drop=True, inplace=True)
```

```
newDf.head()
```

	C	E
0	0.785568	0.655208
1	0.090406	0.017863
2	0.415338	0.643814
3	0.590080	0.627735
4	0.516741	0.041541

## Groupby

The groupby method allows you to group rows of data together and call aggregate functions

```
data = {'Company': ['GOOG', 'GOOG', 'MSFT', 'MSFT', 'FB', 'FB'],
        'Person': ['Sam', 'Charlie', 'Amy', 'Vanessa', 'Carl', 'Sarah'],
        'Sales': [200, 120, 340, 124, 243, 350]}
```

```
df = pd.DataFrame(data)
```

```
df
```

	Company	Person	Sales
0	GOOG	Sam	200
1	GOOG	Charlie	120
2	MSFT	Amy	340
3	MSFT	Vanessa	124
4	FB	Carl	243
5	FB	Sarah	350

Now you can use the `.groupby()` method to group rows together based off of a column name.

For instance let's group based off of Company. This will create a `DataFrameGroupBy` object:

```
df.groupby('Company')
```

```
<pandas.core.groupby.generic.DataFrameGroupBy object at 0x000001918A842B30>
```

You can save this object as a new variable:

```
by_comp = df.groupby("Company")
type(by_comp)
```

```
pandas.core.groupby.generic.DataFrameGroupBy
```

And then call aggregate methods of the object:

```
by_comp.mean(numeric_only=True)
```

	Sales
Company	
FB	296.5
GOOG	160.0
MSFT	232.0

```
df.groupby('Company').mean(numeric_only=True)
```

	Sales
Company	

FB	296.5
GOOG	160.0
MSFT	232.0

```
by_comp.std(numeric_only=True)
```

	Sales
Company	
FB	75.660426
GOOG	56.568542
MSFT	152.735065

```
by_comp.min()
```

	Person	Sales
Company		
FB	Carl	243
GOOG	Charlie	120
MSFT	Amy	124

```
by_comp.max()
```

	Person	Sales
Company		
FB	Sarah	350
GOOG	Sam	200
MSFT	Vanessa	340

```
by_comp.count()
```

	Person	Sales
Company		
FB	2	2
GOOG	2	2
MSFT	2	2

```
by_comp.describe()
```

	Sales							
	count	mean	std	min	25%	50%	75%	max
Company								
FB	2.0	296.5	75.660426	243.0	269.75	296.5	323.25	350.0
GOOG	2.0	160.0	56.568542	120.0	140.00	160.0	180.00	200.0
MSFT	2.0	232.0	152.735065	124.0	178.00	232.0	286.00	340.0